

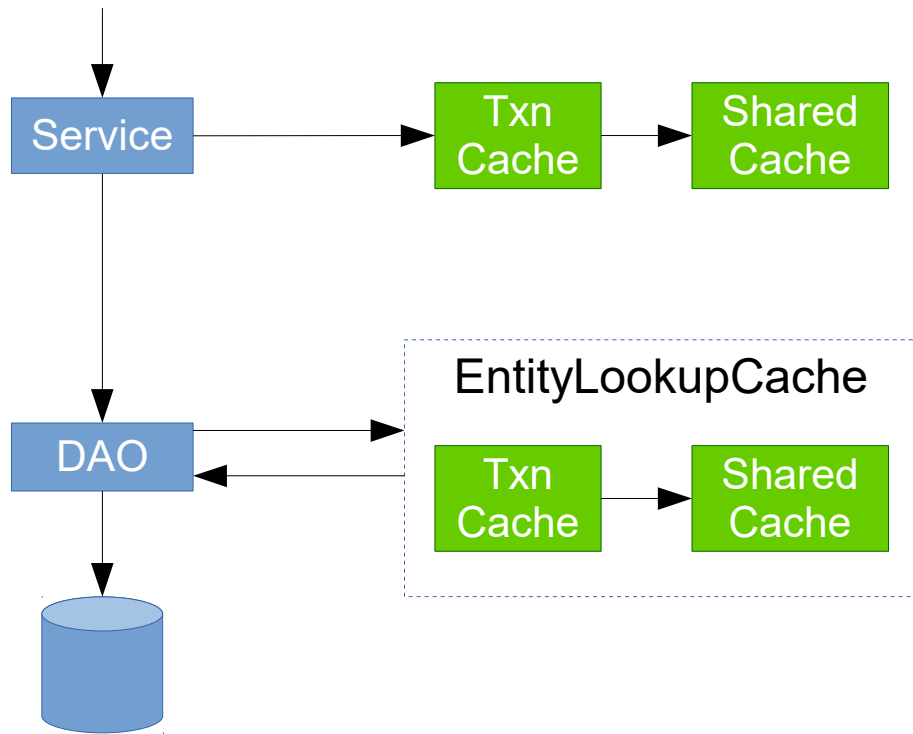
The art of the cache

Axel Faust, Acosix GmbH

Why talk about caches?

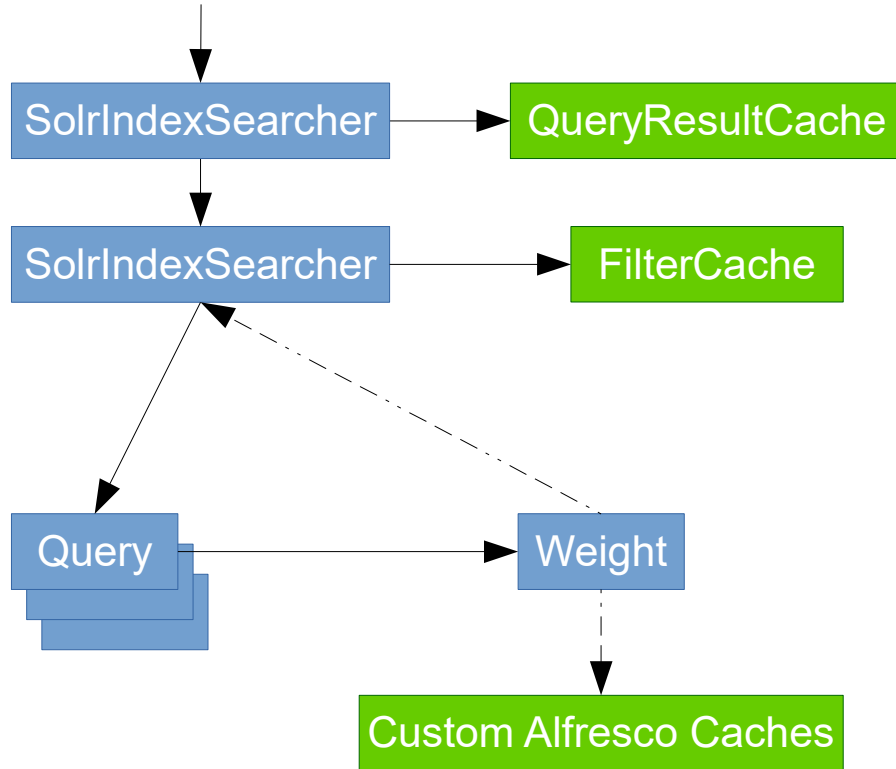
- Everyone wants “performance”
- Little documentation around
 - Configuration options
 - Tuning guides without explanation
- “Art” is proper understanding

Focus: Repository caching



- Multi-layered caching
 - Transaction vs. global
 - Default vs. custom caches
- Technically “pluggable”
 - Declared via config
 - Generic API
- EE-only: distributed data

Focus: SOLR caching



- Multi-layered caching
 - Sub-query recursion
 - + low-level caches
- Easy to understand?
 - Conditional use
 - Indirect relevance
- Technically “pluggable”

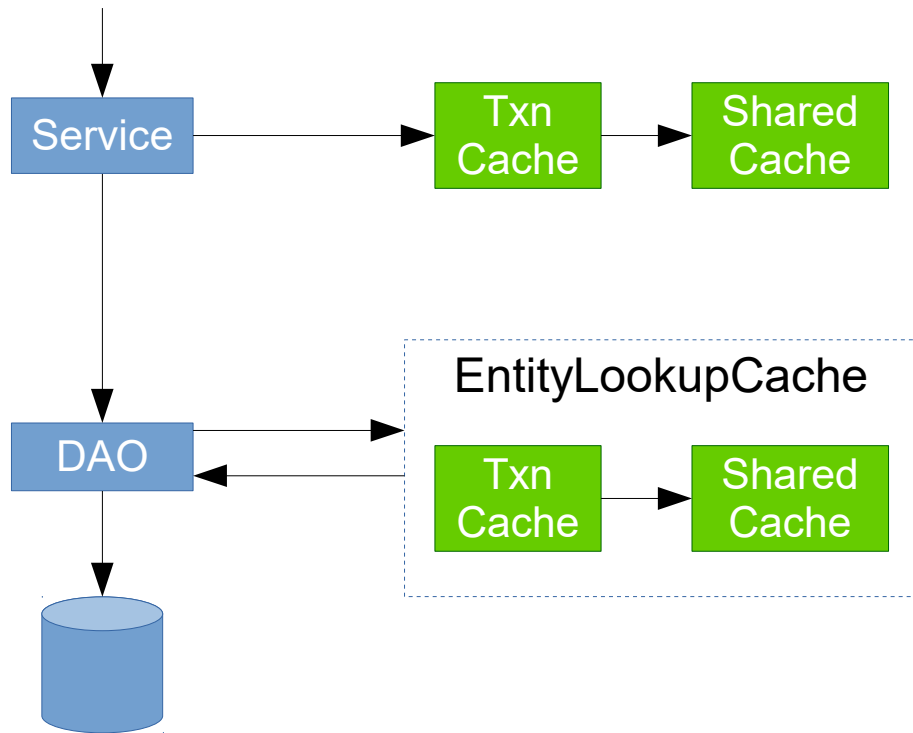
Out of scope

- HTTP caching
 - Client (browser) or HTTP proxy (Apache / nginx)
 - Custom web scripts: use “cache” root object
- Surf caching
 - Model objects (pages, components...)
 - Read-only requests (per Surf request)
- Content caching
 - E.g. CachingContentStore

The background features a cluster of hexagons in light blue, light green, and light yellow. A thick green line forms a large, irregular hexagonal shape that frames the text.


Repository Caches

Components



- SimpleCache (base API)
 - 3 base impl. + Hazelcast (EE)
- TransactionalCache
 - Txn-aware facade
 - Thread-local data until commit
- EntityLookupCache
 - Common DAO logic
 - Callback for specifics

Caching in a cluster (Enterprise)



- Better: “cluster due to caching”
- Mode 1: Fully-distributed
 - 100% Hazelcast
 - Distributed read / write (incl. backups)
- Mode 2: Invalidating
 - Local cache + asynch invalidation
 - Best-effort consistency

Cache definition (Spring beans)

```
<bean name="myAcmeSharedCache" factory-bean="cacheFactory" factory-method="createCache">  
    <constructor-arg value="cache.myAcmeCache"/>  
</bean>
```

```
<bean name="myAcmeCache" class="org.alfresco.repo.cache.TransactionnalCache">  
    <property name="sharedCache" ref="myAcmeSharedCache" />  
    <property name="name" value="com.acme.myCache" />  
    <property name="maxCacheSize" value="${cache.myAcmeCache.tx.maxItems}" />  
    <property name="mutable" value="true" />  
    <property name="disableSharedCache" value="${system.cache.disableMutableSharedCaches}" />  
    <property name="tenantAware" value="false" />  
    <property name="cacheStats" ref="cacheStatistics"/>  
    <property name="cacheStatsEnabled" value="${cache.myAcmeCache.tx.statsEnabled}"/>  
</bean>
```

- Generic pattern – simple to abstract, e.g. via bean factory / emitter

Cache definition (properties)

```
cache.myAcmeCache.tx.maxItems=1000  
cache.myAcmeCache.tx.statsEnabled=${caches.tx.statsEnabled}  
cache.myAcmeCache.maxItems=10000  
cache.myAcmeCache.timeToLiveSeconds=300  
cache.myAcmeCache.maxIdleSeconds=0  
cache.myAcmeCache.eviction-policy=LRU  
cache.myAcmeCache.cluster.type=invalidating  
cache.myAcmeCache.backup-count=1  
cache.myAcmeCache.eviction-percentage=25  
cache.myAcmeCache.merge-policy=hz.ADD_NEW_ENTRY  
cache.myAcmeCache.readBackupData=false
```

- Enterprise: “cluster.type” + below
- Community: “eviction-policy” != NONE => maxItems respected

What caches exist? (default: 52 in 5.2)

| Cache | Configured type | Class | Size | Configured limit | # gets | # hits | Hit % | # misses | Miss % | # evictions |
|--------------------------------------|-------------------|--|--------|------------------|----------|---------|-------|----------|--------|-------------|
| node.aspectsSharedCache | local | org.alfresco.repo.cache.DefaultSimpleCache | 130000 | 130000 | 9387383 | 283982 | 3.0 | 9103401 | 97.0 | 2903509 |
| node.propertiesSharedCache | local | org.alfresco.repo.cache.DefaultSimpleCache | 130000 | 130000 | 9388450 | 285000 | 3.0 | 9103450 | 97.0 | 2903512 |
| nodeOwnerSharedCache | fully-distributed | org.alfresco.repo.cache.DefaultSimpleCache | 40000 | 40000 | 9099324 | 43 | 0.0 | 9099281 | 100.0 | 2993093 |
| node.nodesSharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 23357 | 250000 | 16028609 | 852985 | 5.3 | 15175624 | 94.7 | 6044187 |
| immutableEntitySharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 408 | 50000 | 1400017 | 1397752 | 99.8 | 2265 | 0.2 | 0 |
| contentUriSharedCache | fully-distributed | org.alfresco.repo.cache.DefaultSimpleCache | 68 | 130000 | 170 | 0 | 0.0 | 170 | 100.0 | 0 |
| propertyValueSharedCache | | org.alfresco.repo.cache.DefaultSimpleCache | 44 | | 170 | 20 | 11.8 | 150 | 88.2 | 15 |
| contentDataSharedCache | fully-distributed | org.alfresco.repo.cache.DefaultSimpleCache | 35 | 130000 | 105 | 0 | 0.0 | 105 | 100.0 | 0 |
| node.childByNameSharedCache | local | org.alfresco.repo.cache.DefaultSimpleCache | 20 | 130000 | 611 | 303 | 49.6 | 308 | 50.4 | 0 |
| routingContentStoreSharedCache | local | org.alfresco.repo.cache.DefaultSimpleCache | 13 | 10000 | 53 | 14 | 26.4 | 39 | 73.6 | 0 |
| authoritySharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 10 | 10000 | 44 | 14 | 31.8 | 30 | 68.2 | 0 |
| immutableSingletonSharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 9 | 12000 | 182 | 161 | 88.5 | 21 | 11.5 | 0 |
| propertyClassSharedCache | | org.alfresco.repo.cache.DefaultSimpleCache | 8 | | 34 | 14 | 41.2 | 20 | 58.8 | 0 |
| propertyUniqueContextSharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 7 | 10000 | 33 | 6 | 18.2 | 27 | 81.8 | 0 |
| aclEntitySharedCache | fully-distributed | org.alfresco.repo.cache.DefaultSimpleCache | 7 | 50000 | 44 | 23 | 52.3 | 21 | 47.7 | 0 |
| node.rootNodesSharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 6 | 1000 | 30581 | 30563 | 99.9 | 18 | 0.1 | 0 |
| authorityToChildAuthoritySharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 5 | 40000 | 20 | 5 | 25.0 | 15 | 75.0 | 0 |
| openCMISRegistrySharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 2 | 500 | 4 | 4 | 100.0 | 0 | 0.0 | 0 |
| personSharedCache | fully-distributed | org.alfresco.repo.cache.DefaultSimpleCache | 2 | 1000 | 31 | 22 | 71.0 | 9 | 29.0 | 0 |
| ticketsCache | fully-distributed | org.alfresco.repo.cache.DefaultSimpleCache | 2 | 1000 | 29 | 29 | 100.0 | 0 | 0.0 | 0 |

Impact of caches

- Searches: `node.*Cache(s)`
 - Example SOLR benchmark: 312 ms (1st run) vs. 135 ms (2nd run)
- Memory footprint: `immutableEntitySharedCache`
 - Namespace, QName, mimetypes et al
 - Saves e.g. <~ 350 byte per QName re-use
- User-group lookups: `userToAuthoritySharedCache`
 - Avoid recursive DB lookup over parent group tree

Using custom caches

- DAO: stick to EntityLookupCache
- Checklist questionnaire
 - Is caching really a necessity?
 - Miss cost vs. memory cost
 - Are types of keys / values suitable?
 - How to detect / avoid stale data?
 - Versioned key vs. invalidation by behaviour / scheduled action



Repository Caches Tuning

Understanding the issue

“Tuning the caches in a wrong way may lead to out of memory issues.”

- Sources of information
 - Logs: “Transactional update cache '[name]' is full ([limit])”
 - Monitoring: JVM heap / garbage collection
 - Code: cache usage patterns
 - Tuning guides: be wary of source / age
- No cache insights in default Alfresco
 - => OOTBee Support Tools

Understanding the issue

| Cache | Configured type | Class | Size | Configured limit | # gets | # hits | Hit % | # misses | Miss % | # evictions |
|--------------------------------------|-------------------|--|--------|------------------|----------|---------|-------|----------|--------|-------------|
| node.aspectsSharedCache | local | org.alfresco.repo.cache.DefaultSimpleCache | 130000 | 130000 | 9387383 | 283982 | 3.0 | 9103401 | 97.0 | 2903509 |
| node.propertiesSharedCache | local | org.alfresco.repo.cache.DefaultSimpleCache | 130000 | 130000 | 9388450 | 285000 | 3.0 | 9103450 | 97.0 | 2903512 |
| nodeOwnerSharedCache | fully-distributed | org.alfresco.repo.cache.DefaultSimpleCache | 40000 | 40000 | 9099324 | 43 | 0.0 | 9099281 | 100.0 | 2993093 |
| node.nodesSharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 23357 | 250000 | 16028609 | 852985 | 5.3 | 15175624 | 94.7 | 6044187 |
| immutableEntitySharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 408 | 50000 | 1400017 | 1397752 | 99.8 | 2265 | 0.2 | 0 |
| contentUriSharedCache | fully-distributed | org.alfresco.repo.cache.DefaultSimpleCache | 68 | 130000 | 170 | 0 | 0.0 | 170 | 100.0 | 0 |
| propertyValueSharedCache | | org.alfresco.repo.cache.DefaultSimpleCache | 44 | | 170 | 20 | 11.8 | 150 | 88.2 | 15 |
| contentDataSharedCache | fully-distributed | org.alfresco.repo.cache.DefaultSimpleCache | 35 | 130000 | 105 | 0 | 0.0 | 105 | 100.0 | 0 |
| node.childByNameSharedCache | local | org.alfresco.repo.cache.DefaultSimpleCache | 20 | 130000 | 611 | 303 | 49.6 | 308 | 50.4 | 0 |
| routingContentStoreSharedCache | local | org.alfresco.repo.cache.DefaultSimpleCache | 13 | 10000 | 53 | 14 | 26.4 | 39 | 73.6 | 0 |
| authoritySharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 10 | 10000 | 44 | 14 | 31.8 | 30 | 68.2 | 0 |
| immutableSingletonSharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 9 | 12000 | 182 | 161 | 88.5 | 21 | 11.5 | 0 |
| propertyClassSharedCache | | org.alfresco.repo.cache.DefaultSimpleCache | 8 | | 34 | 14 | 41.2 | 20 | 58.8 | 0 |
| propertyUniqueContextSharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 7 | 10000 | 33 | 6 | 18.2 | 27 | 81.8 | 0 |
| aclEntitySharedCache | fully-distributed | org.alfresco.repo.cache.DefaultSimpleCache | 7 | 50000 | 44 | 23 | 52.3 | 21 | 47.7 | 0 |
| node.rootNodesSharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 6 | 1000 | 30581 | 30563 | 99.9 | 18 | 0.1 | 0 |
| authorityToChildAuthoritySharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 5 | 40000 | 20 | 5 | 25.0 | 15 | 75.0 | 0 |
| openCMISRegistrySharedCache | invalidating | org.alfresco.repo.cache.DefaultSimpleCache | 2 | 500 | 4 | 4 | 100.0 | 0 | 0.0 | 0 |
| personSharedCache | fully-distributed | org.alfresco.repo.cache.DefaultSimpleCache | 2 | 1000 | 31 | 22 | 71.0 | 9 | 29.0 | 0 |
| ticketsCache | fully-distributed | org.alfresco.repo.cache.DefaultSimpleCache | 2 | 1000 | 29 | 29 | 100.0 | 0 | 0.0 | 0 |

“Transactional update cache is full”

- Cause: too many read/updated or removed entries
 - put() trigger: $\#uniqueKeys(read + update) \geq \#max$
 - remove() trigger: $\#keys(remove) \geq \#max$
- Quick & dirty: `cache.XXX.tx.maxItems`
- Ideal: analyse + fix use case
 - Smaller batch sizes / scope of work
 - Single txn recursion => multi-txn batch processing


Shared cache optimisations

- Sizing considerations
 - Fit in JVM heap / avoid GC pressure (use adaptive GC!)
 - Support X% of total DB entries
 - Support Y minutes worth of K-times user access repetition
 - Indicator 1: evictions per Y minutes \leq cache size
 - Indicator 2: hit count per Y minutes $\geq K * \text{cache size}$
- Option: pre-load expected data sets
 - Asynch txn post-processing / scheduled process
 - Based on business metadata, access or query logs

Estimating memory consumption

- Recommendation: load realistic data sets + analyse heap dump
 - Too many factors for generic formulas
 - Highly dependent on model / user data
- Personal experience
 - node.nodesSharedCache: $\#entries * \sim 700 \text{ byte}$
 - node.propertiesSharedCache: $\#entries * \text{avg}(\#propsPerNode) * \sim 100 \text{ byte}$
 - node.aspectsSharedCache: $\#entries * (\sim 150 \text{ byte} + \text{avg}(\#aspectsPerNode) * 32 \text{ byte})$
 - content*SharedCache: $\#entries * \sim 150 \text{ byte}$
 - General overhead: $\#entries * \sim 125 \text{ byte}$

Extreme: custom implementation

- Example: -backed module
 - Off-heap data storage = less GC pressure
 - Distributed cache, compute and messaging
- Approach
 - Implement CacheFactory / SimpleCache
 - Deal with “weird” Alfresco cache use-cases
 - E.g. immutableSingletonSharedCache, globalConfigSharedCache...
 - Optional: use original cache properties as default


Annoying cache bits

- TTL on `node.nodesSharedCache` unnecessary
 - Default 300s quite aggressive
 - Significant for bulk-fetch, e.g. during search
- Obsolete Hazelcast version (2.4 vs. 3.8)
- Non-standard cache: `parentAssocsCache`
 - Not measurable
 - Limited configurability – only `maxSize` + “`limitFactor`”

The background features a pattern of hexagons in light blue, light green, and light yellow. A thick green line is drawn across the image, forming a large, irregular shape that frames the text.

SOLR Caches

Caches in use



– SOLR default caches

- queryResultCache
- filterCache
- documentCache (highlight/tracking)
- fieldValueCache (facets)

– Alfresco default caches

- denied/authorityCache (ACL checks)
- reader/ownerCache (ACL checks - indirect)
- pathCache

Cache relevance

- Unscientific comparisons (4m docs, SSD)
 - PROPERTIES:"<qname>": ~35ms vs. 1ms
 - PATH:"<path>": ~1s vs. 1ms
 - Facets of unique-valued field: ~750ms vs ~150ms
- Sort by score prevents some cache use
- Typical cache entry bytes: $\sim \text{\#index docs} / 8 (+\text{key})$
 - QueryResultCache: $\sim \text{\#resultsInSlice} * 4$

Understanding “auto-warm”

- Cache data “refresh”
- Default limits quite low
 - 4 authority and 32 path sub-queries, 32 filter queries
 - 4 complete result sets
- Option: reject or off-load complex / costly entries
 - Interface for custom code: CacheRegenerator

Cache tuning

- Typical focus: high hit ratio – low warmupTime
 - Problem: layering + arbitrary queries skew metrics
- General priority: I/O caches + throughput
- Different system = different approach
 - No warmup if update frequency high
 - Increased cache size if used queries granular / re-usable
 - Custom regenerator for expensive queries – or no warmup

The background features a pattern of hexagons in light blue, light green, and light yellow. A thick green line forms a large, irregular hexagonal shape that frames the text. The text is centered within this green frame.

“The art of the cache”

Levels

- Basic
 - Understand default behavior
 - Able to judge metrics + tune accordingly (filter out the noise / bad advice)
- Advanced
 - Build custom solutions utilising custom caches
 - Implement tweaks / hooks
- ...go “wild” beyond that



@ReluctantBird83

axel.faust@acosix.org