

Solr lab for sharding experiments

Roxana Angheluta (Xenit Solutions)

Axel Faust (Acosix GmbH)

Motivation (problem)

Requirements

- Large archives with >50mil documents
- No sites
- Search speed <3s
- Particular use case: most of the queries are targeted to a specific client id

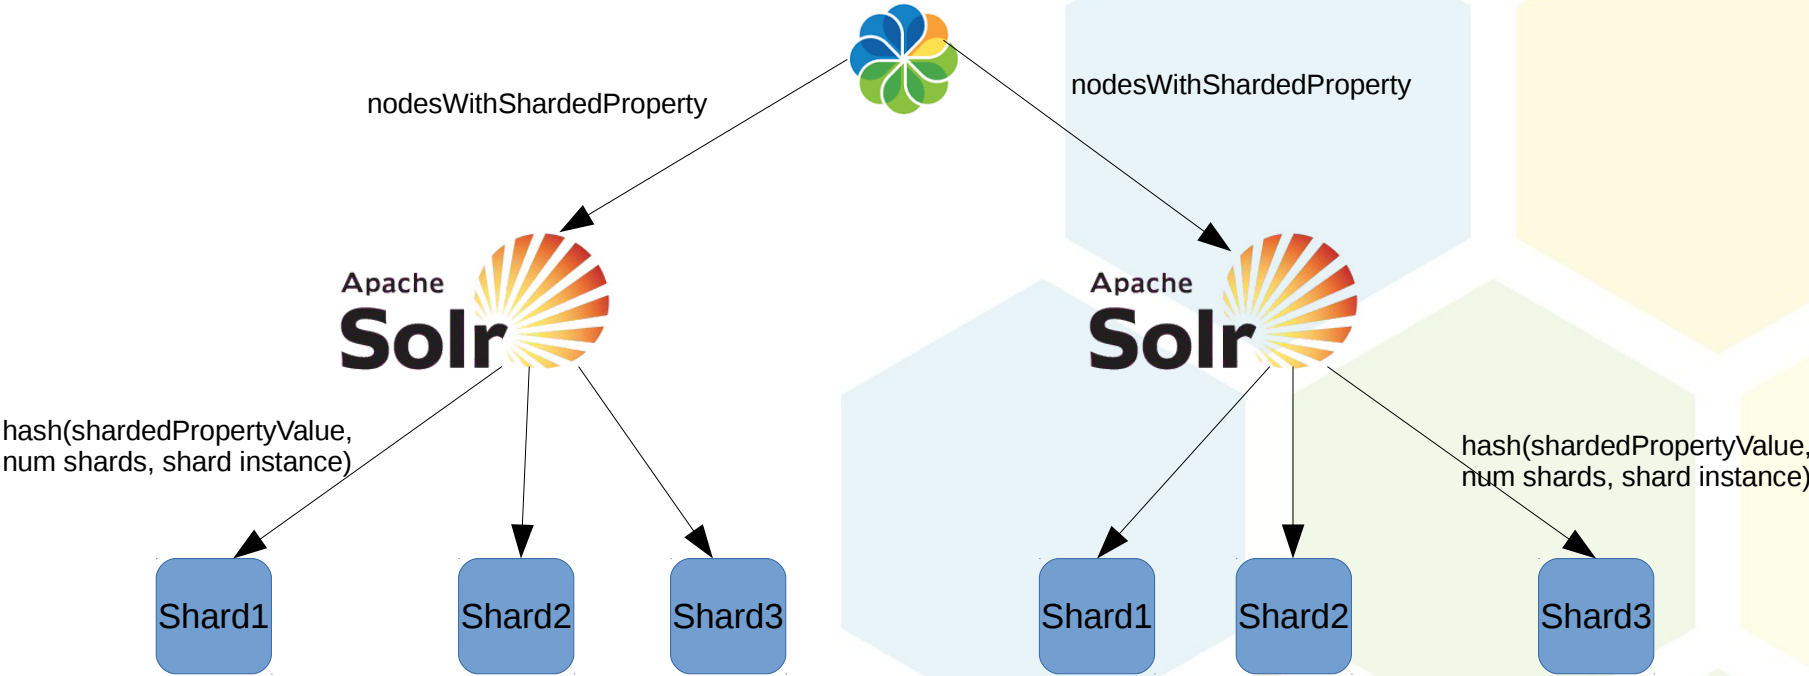


Solution

- Implement a new sharding method: by property (now existing in Alfresco 5.2)
- Create a “solr lab”
- Define relevant parameters, metrics
- Design relevant experiments
- Execute experiments
- Monitor
- Get results
- Draw conclusions



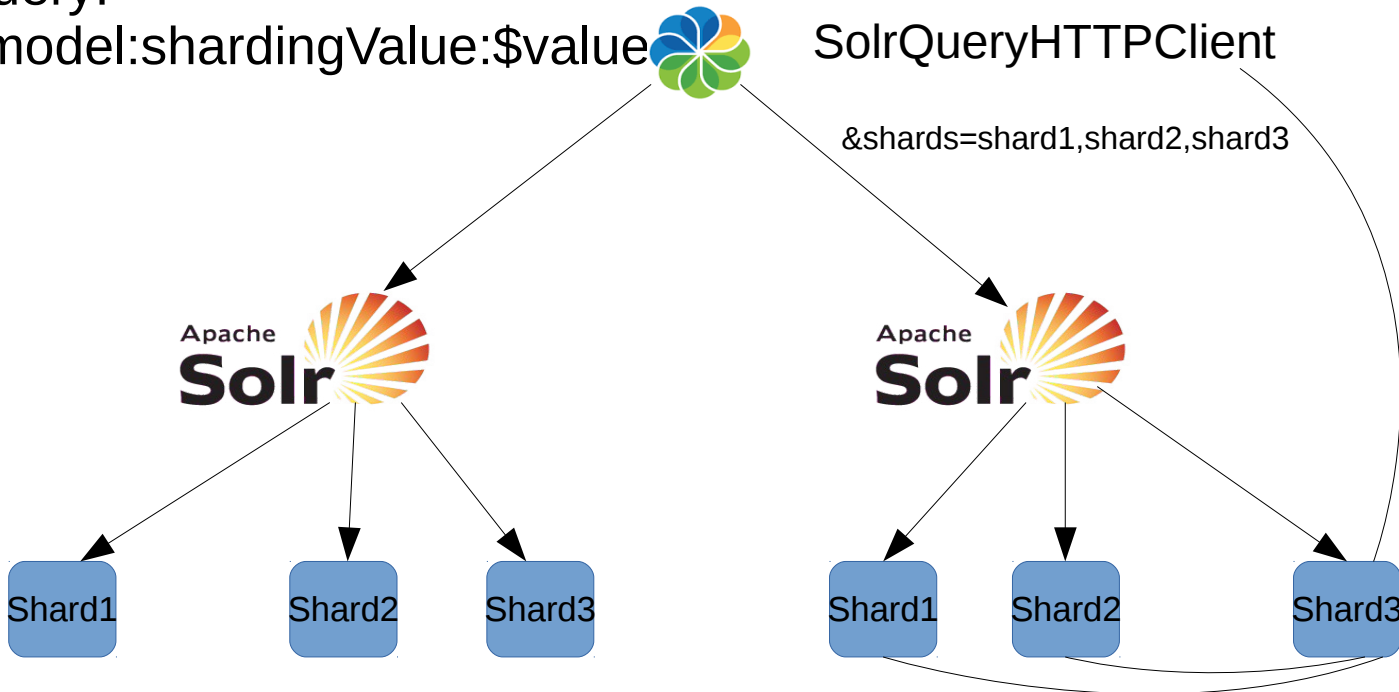
Sharding by property - Solr side



Sharding by property - Alfresco side

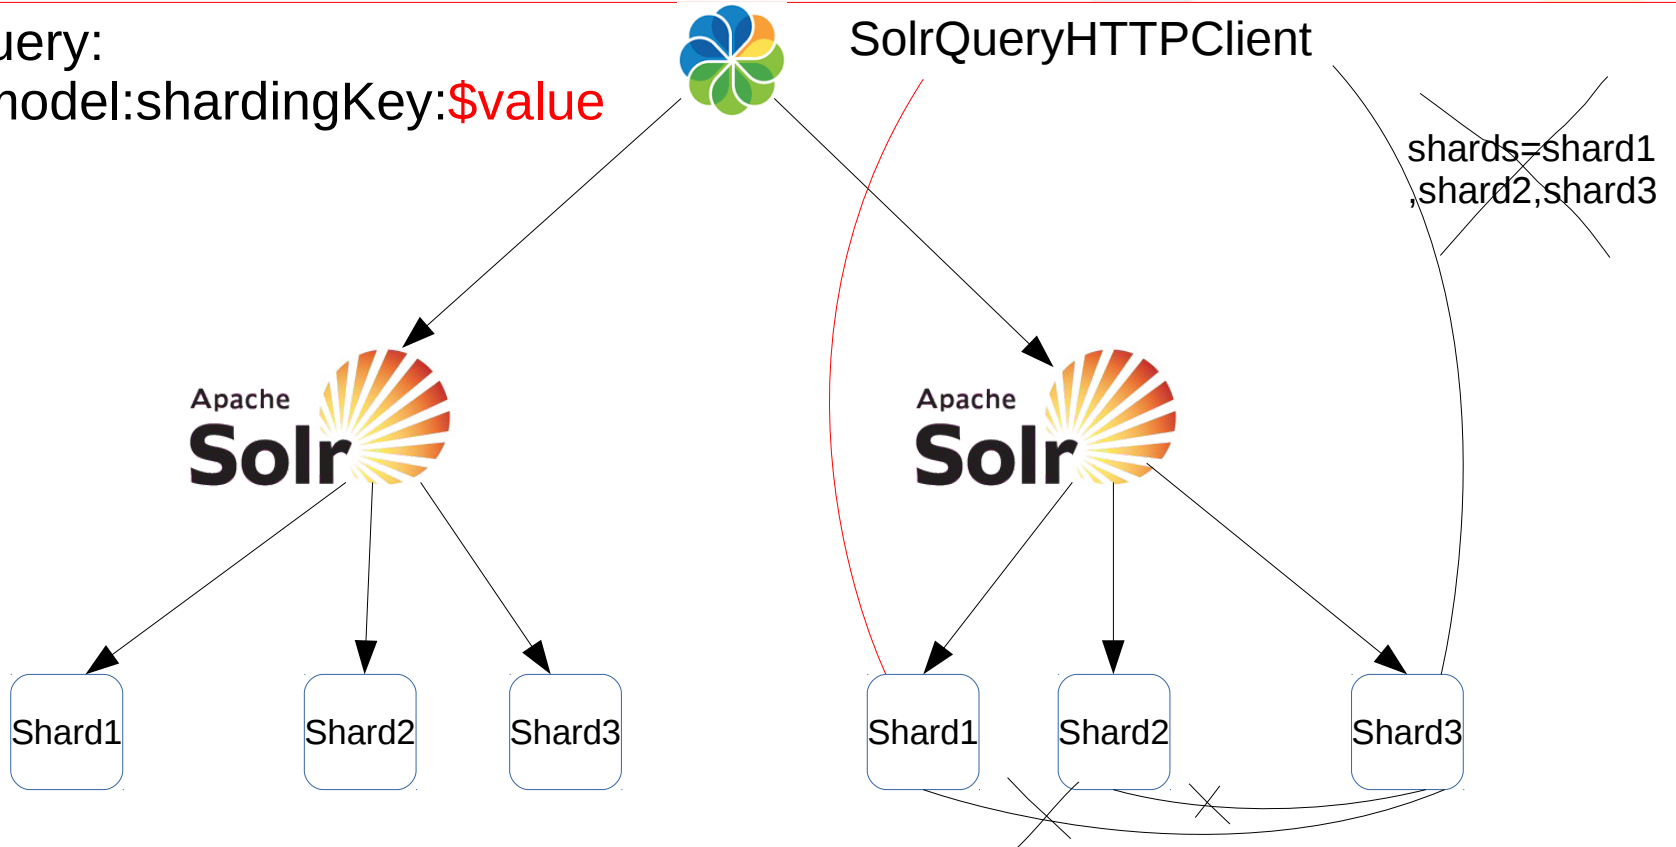
- Query:

=model:shardingValue:\$value



Sharding by property - Alfresco side

- Query:
=model:shardingKey:\$value



Solr lab

- Requirements
 - Easy to start
 - Easy to replicate
 - Easy to change parameters
 - Easy to resume
 - “Monitorable”



Solr lab

docker-compose.yml


```
core1:
  image:
  volumes:
  hostname:
  ports:
  environment:
solr1:
...
solr2:
...
solr3:
...
postgres:
....
libreoffice
```



Alfresco



Solr shard1



PostgreSQL

Database

Solr lab

core1:

image: image: hub.xenit.eu/alfresco-ethias-sharded-by-property:build-3

volumes:

- alfresco10milnoacls:/opt/alfresco/alf_data
- ./license:/opt/alfresco/license

hostname:

ports:

environment:

.....

- DB_NAME=alfresco
- DB_USERNAME=alfresco
- DB_PASSWORD=admin
- DB_URL=jdbc:postgresql://\${NET_PRIVATE_IP}:7432/alfresco
- SOLR_HOST=solr1
- SOLR_PORT=8081
- ALFRESCO_HOST=core1
- ALFRESCO_PORT=8080
- JAVA_XMS=6144M
- JAVA_XMX=6144M
- ENABLE_CLUSTERING=true
- DYNAMIC_SHARD_REGISTRATION=true
- SOLR_SSL=none
- SERVICE_8080_NAME=alfresco
- SERVICE_8080_TAGS=proxy-http
- LIBREOFFICE_HOST=libreoffice
- LIBREOFFICE_PORT=8997

- DEBUG=false
- JMX_ENABLED=true
- JMX_RMI_HOST=xxx.xxx.xxx.xxx
- GLOBAL_shardedProperty.qname={http://www.ethias.be/model/content}shardingKey
- GLOBAL_cache.node.nodesSharedCache.maxItems=1250000
- GLOBAL_cache.node.nodesSharedCache.timeToLiveSeconds=3600
- GLOBAL_cache.node.aspectsSharedCache.maxItems=650000
- GLOBAL_cache.node.aspectsSharedCache.timeToLiveSeconds=3600
- GLOBAL_cache.node.propertiesSharedCache.maxItems=650000
- GLOBAL_cache.node.propertiesSharedCache.timeToLiveSeconds=3600

Solr lab

solr1:

image: hub.xenit.eu/solr-sharded-by-property:5.1

volumes:

- solr10milnoacsl1index:/opt/alfresco/alf_data
- solr10milnoacsl1conf:/opt/alfresco/solr4
- ./log4j-solr.properties:/opt/alfresco/solr4/log4j-solr.properties

ports:

hostname: solr1

environment:

- ALFRESCO_HOST=core1
- ALFRESCO_PORT=8080
- SOLR_HOST=solr1
- SOLR_PORT=8081
- JAVA_XMS=7168M
- JAVA_XMX=7168M
- SHARDING=true
- NUM_SHARDS=3
- NUM_NODES=3
- NODE_INSTANCE=1
- TEMPLATE=rerank
- SHARD_IDS=0
- ALFRESCO_SSL=none

- ARCHIVE_ENABLE_ALFRESCO_TRACKING=false
- ARCHIVE_INDEX_CONTENT=false
- ALFRESCO_INDEX_CONTENT=false
- MAX_HTTP_HEADER_SIZE=65536
- JMX_ENABLED=true
- JMX_RMI_HOST=136.243.138.174

Experiments - dimensions

Size of repository: 1mil, 10mil, more

Sharding method: acl, property

Acls: with / without artificial acls for a more fair distribution of documents in the acl sharding method

Number of shards

Types of queries: simple, complex (booleans + facets)

Garbage collector used

Concurrency

With / without loading in the same time



Experiments - metrics

- Important to measure
 - Response times
 - Memory usage
 - Load
 - Caches utilization
 - Database performance



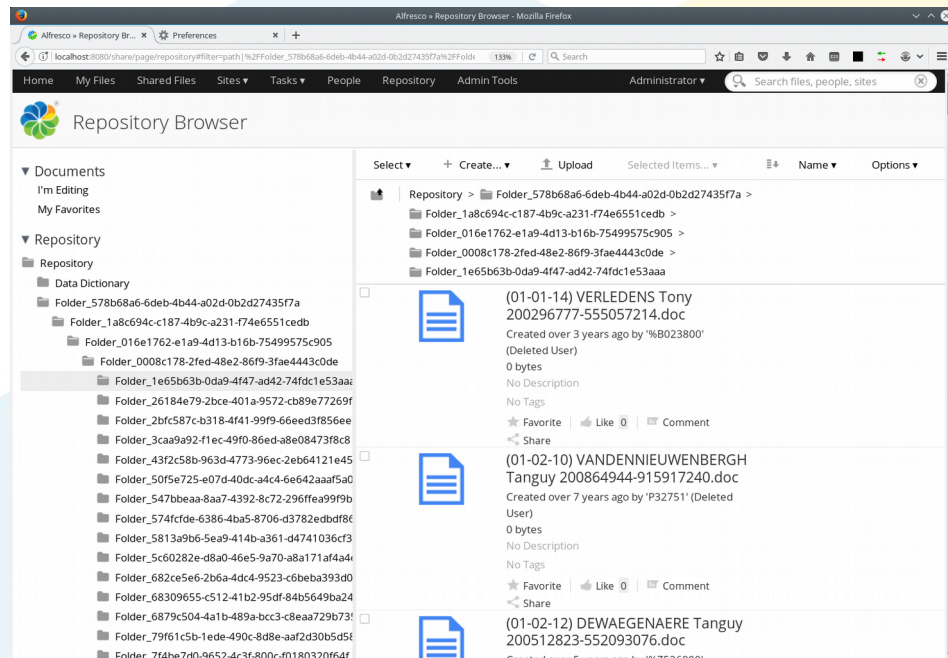
Experiments - hardware

- Started experiments at Amazon, with AuroraDB as database
- Moved to Hetzner dedicated servers, tried multiple configurations, stabilized to:
https://www.hetzner.de/nl/hosting/produkte_rootserver/px121ssd
 - Intel(R) Xeon(R) CPU E5-1650 v3 @ 3.50GHz, 12 cores
 - 264GB RAM
 - 2TB SSD disks

		Database	Index 1 shard
1 mil	1 shard	8.7G	5.1G
	3 shard		1.8G
10 mil	1 shard	78G	51G
	3 shard		17G
	11 shard		4.8G
100 mil	3 shard	1.3T	164G


Experiments - loading documents

- No content
- Custom model with >10 metadata fields, values randomly generated based on constraints or data type
- Folder structure with / without artificial acls
- Load speed at 12 threads:
 - Approx 1.3 mil docs / hour
- Solr indexing real time



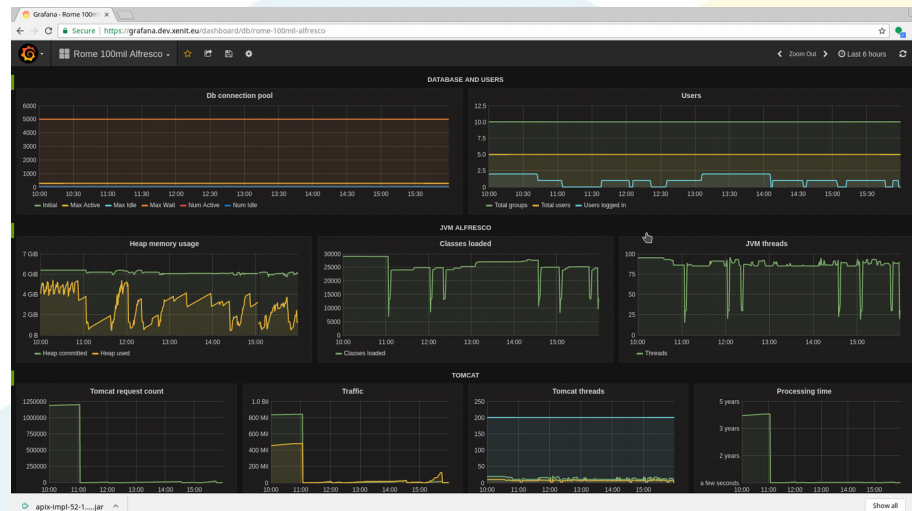
Experiments – problems encountered

- Ran out of inodes on solr' partition – needed to reformat the partition with small block size
- Db tuning – max connections
- Memory tuning on db, alfresco + solr size
- Increased max header size in tomcat



Monitoring

- Jmeter
- Jvisualvm
- Grafana stack
 - Jmxtrans to collect jmx metrics
 - Grafana container



Results

Size	Shards	Context	Method	Artificial ACLs	Threads	Queries per thread	Run	Avg (ms)	Median (ms)	90% (ms)	99% (ms)	Max (ms)	Throughput/s
10mil	3	regular	acl	yes	16	150	#1	456	375	878	1433	3937	33,1
							#2	181	165	292	437	573	77,7
			property	yes	16	150 (after write tests)	#1	292	237	539	890	3496	51,1
							#2	412	324	859	1486	2088	35,0
		regular				100 (after write tests)	#1	88	85	119	167	243	142
							#2	336	219	693	1152	4367	43,2
			acl	no	16	150	#1	465	342	938	3174	4605	30,8
							#2	111	95	162	375	3102	103,3
		with facets	acl	no	16	150	#1	596	490	1165	2206	4378	25,8
							#2	115	74	251	517	646	120,2
			property	no	16	150	#1	404	298	734	2849	4734	36,5
							#2	88	81	116	329	422	143,2
		with facets	acl	yes	16	120	#1	919	767	1554	3498	10588	16,3
							#2	937	778	1645	3190	9628	16,3
							#1	385	364	540	764	1165	38,3
							#2	387	365	545	852	1054	38,5
		regular			8	240	#1	466	367	800	1871	9207	16,4
							#2	336	321	442	580	766	22,5

Conclusions

- Search slow?
- Primary culprit: DB access / cache
 - Expensive bulk load
 - Cache “sabotage”: nodesSharedCache TTL (used OOTBee Support Tools for insights)



Conclusions

- Search slow?
- Primary culprit: DB access / cache
 - Expensive bulk load
 - Cache “sabotage”: nodesSharedCache TTL (used OOTBee Support Tools for insights)
- Secondary culprit: Repository-tier ACL checking
 - Misleading: PermissionEvaluationMode.NONE
 - Low default size of caches for readers/readersDenied

Conclusions

- Natural best performer: property sharding + shard-targeted queries
 - ~20% (+-10%) better general performance
 - Improved scaling for parallel requests



Further work

- In progress: 100 million documents
 - Extended data model
 - Realistic value distribution for sharding property
- Planned: customer reference setup
 - Ideally: 100% live data clone, augmented via generation
 - Oracle instead of PostgreSQL



Speaker contacts

roxana.angheluta@xenit.eu

axel.faust@acosix.de

